


What boring AI looks like

Six plays for a company where AI is in production, measured, and trusted exactly as far as it's reliable.

Start with your weakest play. You're only as boring as your weakest link.

THE FOUR STAGES



PUTS INTO PRACTICE

- Bo** [AI Should Be Boring](#) **Rt** [Risk-Tiered AI Adoption](#) **Ff** [Fix First, AI Last](#) **Ct** [Calibrated Trust](#) **Kh** [Keep It Human](#)

01
Rt RISK-TIERED AI ADOPTION

Sequence by risk

Build the muscle where mistakes are cheap.

Catching AI's mistakes is a skill. Practice it inside first.

THREE TIERS, LIGHT REVIEW TO HEAVY

- Internal productivity**
→ a person sees every output
- Operations**
→ mistakes reach colleagues
- Customer-facing**
→ mistakes reach customers

YOU'LL KNOW IT'S WORKING WHEN

- ✓ Ideas are sorted by risk first, then by return within each tier.
- ✓ Risk means the data a tool touches and what it can do on its own.
- ✓ Internal tools are in daily use, and their lessons shape what comes next.

NOT THERE YET WHEN

- Ideas move when someone senior gets excited.
- The first big bet faces customers.

BRIDGE THE GAP

- Sort ideas into the three tiers, then rank by return within each.
- Put one tool into daily use where mistakes are cheap and easy to spot.

02
Ff FIX FIRST, AI LAST

Fix the work first

Don't automate the jam.

AI on a broken process just makes the mess faster.

THE FOUR QUESTIONS, IN ORDER

- Chasing, waiting or redoing?**
→ Fix the process
- Same steps every time?**
→ Automate it with a rule
- Mostly reading or writing?**
→ AI drafts, a person checks
- Anything else?**
→ Keep it human

YOU'LL KNOW IT'S WORKING WHEN

- ✓ Processes are mapped step by step, with chasing, waiting and redoing marked.
- ✓ Causes get fixed first, then simple rules. AI gets the reading and writing left.
- ✓ Data sits in systems approved tools can reach, with an owner keeping it accurate.

NOT THERE YET WHEN

- A tool is proposed before anyone knows the cause.
- The real process lives in a few people's heads.

BRIDGE THE GAP

- Map one process this week. Circle the chasing, waiting and redoing.
- Put each stuck step through the four questions. The first yes decides.

03
Rt RISK-TIERED AI ADOPTION

Govern in proportion

One door, two speeds.

Heavy review for everything teaches people to route around it.

TWO SPEEDS

If the request is	IT GETS
Low risk	A quick yes
On customer data, or acting alone	Security review first

YOU'LL KNOW IT'S WORKING WHEN

- ✓ One place to ask: low-risk requests get a quick yes, riskier ones a closer look.
- ✓ Security reviews tools that touch customer data before the build, on a set turnaround.
- ✓ A current register lists every AI tool, its risk tier, data and owner.

NOT THERE YET WHEN

- People use tools nobody approved.
- Security first sees a tool after it's built.

BRIDGE THE GAP

- Open one place to ask, with a promised turnaround.
- Start a one-page register: tool, purpose, data, tier, owner. Review it monthly.

04

Ct CALIBRATED TRUST

Trust the track record

Use AI where checking is cheaper than doing.

AI sounds just as sure when it's wrong. Fluency isn't truth.

THE TRUST MAP

If a mistake is	EASY TO CHECK	HARD TO CHECK
Minor	Delegate freely	Ideas only
Costly	Use, then verify	Keep it human

YOU'LL KNOW IT'S WORKING WHEN

- ✓ Each use is placed by the cost of a mistake and the ease of checking.
- ✓ An expert checks the output, as a step that can't be skipped.
- ✓ Tools are tested on real examples before launch, and sampled after.

NOT THERE YET WHEN

- Output goes out because it reads well.
- Errors surface as complaints, not as a number.

BRIDGE THE GAP

- Put your five most common AI uses on the trust map.
- Have an expert grade 20 outputs: right, fixable or wrong. Repeat monthly.

05

Kh KEEP IT HUMAN

Keep the judgment

Hand off the hauling, keep the reins.

Judgment is a muscle. Every call you just approve is a skipped rep.

BE THE CENTAUR

The work	CENTAUR	REVERSE CENTAUR
Thinking	Augmented	Outsourced
Doing	Outsourced	Kept
Judging	Kept	Outsourced

YOU'LL KNOW IT'S WORKING WHEN

- ✓ Which calls stay human is agreed in advance, and a named person makes each one.
- ✓ Training is on real tasks: when to use AI, how to check it, what to keep.
- ✓ Named champions have time set aside, share what works and report what breaks.

NOT THERE YET WHEN

- Signing off means approving whatever the tool suggested.
- Training means one session on features.

BRIDGE THE GAP

- Before the work starts, name who makes the final call.
- Before anything goes out, ask: could I defend this without "the AI suggested it"?

06

Bo AI SHOULD BE BORING

Prove it with a number

Agree the number before you build.

A demo shows AI at its best. A number shows an ordinary Tuesday.

EVERY PILOT'S PATH

Baseline → Target → Decision date
→ Scale, change or stop

YOU'LL KNOW IT'S WORKING WHEN

- ✓ A baseline and a target are agreed before the build, with whoever judges success.
- ✓ The number is tracked against the baseline on a schedule, by a named owner.
- ✓ Every pilot has a decision date: scale, change or stop, on the number.

NOT THERE YET WHEN

- Success is reported in logins and licenses.
- Pilots run on because nobody decided to stop them.

BRIDGE THE GAP

- Before the next project, pick one number it should move, and measure it today.
- Give every pilot a decision date. Decide with the number in front of you.

Responsible AI guardrails

Six rules for every tier.

- Rt** Anything acting on customer data, or sending it outside, gets security review before the build.
- Ct** Anything that sends, pays, deletes or publishes gets a person in the loop or a hard limit.
- Rt** Every AI tool has a named owner and a line in the register.
- Kh** A named person makes every call that matters, and can defend it without the AI.
- Ct** When AI decides about people, check its errors group by group: an average can hide who it fails.
- Kh** Say where AI did the work: it's a tool, not a secret.

The AI champion network

Champions are sensors, not cheerleaders.

HOW THE NEWS TRAVELS

What breaks

Team → Champion → Tool owner

What works

Team → Champion → Every team

- One named champion per team, with a few hours a month.
- Every month, two questions: what's working that others should copy? What's breaking that someone should fix?
- They teach on real tasks, not features.

Where do you stand? Score all six plays and get your three biggest gaps. georgeandersen.net/tools/how-boring-is-your-ai