

What boring AI looks like

A two-page guide to getting AI out of pilots and into everyday work, safely.

Why boring?

The most useful technologies end up boring. Nobody gets excited about electricity or spreadsheets; people just rely on them. Most AI isn't there yet: it impresses in demos and stalls in pilots. Boring AI is AI you can depend on:

In production

part of everyday work, not a pilot or a demo

Fit to the work

added after the process is fixed

Measured

judged by a number agreed up front

Trusted as far as it's reliable

checked as much as a mistake would cost

Start here: the first 90 days

1 WEEKS 1-2

Set ground rules: what may go into AI tools, and one place to ask before trying a new one.

Owner: AI lead and security

2 WEEKS 3-4

List every AI tool in use: its owner, its data, its risk.

Owner: AI lead

3 MONTH 2

Pick one internal task where mistakes are cheap. Fix the process, measure the start, then add one tool.

Owner: A team lead

4 MONTH 3

Have an expert grade 20 outputs. Then scale it, change it or stop.

Owner: Sponsor

WHO'S WHO **Sponsor:** an executive who funds the work and decides on pilots. **AI lead:** runs the tool list and the place to ask.

Security: reviews anything touching customer data.

The six plays

Each shows what good looks like, a warning sign and a first move.

01 Sequence by risk

Build skill where mistakes are cheap.

THREE TIERS, LIGHT REVIEW TO HEAVY

- 1 **Help with people's own work**
→ a person sees every output
- 2 **AI inside a business process**
→ mistakes reach colleagues
- 3 **Anything customers see**
→ mistakes reach customers

YOU'LL KNOW IT'S WORKING WHEN

- ✓ Ideas are sorted by risk first, then by return within each tier.
- ✓ Risk means the data a tool touches and what it can do on its own.
- ✓ Internal tools are in daily use, and teach what to try next.

WATCH OUT FOR

- The first big AI bet faces customers.

FIRST MOVE

- Sort your AI ideas into the three tiers, then rank by return within each.

02 Fix the work first

Don't automate the jam.

AI added to a broken process just makes the mess faster.

THE FOUR QUESTIONS, IN ORDER

- 1 **Chasing, waiting or redoing?**
→ Fix the process
- 2 **Same steps every time?**
→ Automate it with a rule
- 3 **Mostly reading or writing?**
→ AI drafts, a person checks
- 4 **Anything else?**
→ Keep it human

YOU'LL KNOW IT'S WORKING WHEN

- ✓ Processes are mapped step by step, with chasing, waiting and redoing marked.
- ✓ Causes are fixed first, then rules; AI gets the reading and writing left.
- ✓ Data is in systems approved tools can reach, and someone keeps it accurate.

WATCH OUT FOR

- A new tool is proposed before anyone knows why the work is slow.

FIRST MOVE

- Put each slow step through the four questions. The first yes decides.

03 Govern in proportion

One front door, two speeds.

Heavy review for everything teaches people to go around it.

TWO SPEEDS

If the request is

IT GETS

Low risk

A quick yes

On customer data, or acting alone

Security review first

YOU'LL KNOW IT'S WORKING WHEN

- ✓ One place to ask: a quick yes for low risk, a closer look for the rest.
- ✓ Security reviews customer-data tools before they're built, on a set turnaround.
- ✓ A current one-page list shows every AI tool's risk, data and owner.

WATCH OUT FOR

- People use AI tools nobody approved.

FIRST MOVE

- Promise a turnaround for requests, and keep it.

04 Trust the track record

Use AI where checking is cheaper than doing.

AI sounds just as sure when it's wrong.

HOW FAR TO RELY ON IT

If a mistake is	EASY TO CHECK	HARD TO CHECK
Minor	Let it run	Ideas only
Costly	Use, then verify	Keep it human

YOU'LL KNOW IT'S WORKING WHEN

- ✓ Each use is placed by the cost of a mistake and the ease of checking.
- ✓ An expert checks the output, as a step that can't be skipped.
- ✓ Tools are tested on real examples before launch, and sampled after.

WATCH OUT FOR

- Output goes out because it reads well.

FIRST MOVE

- Place your five most common AI uses on this grid.

05 Keep people in charge

AI does the work. People make the call.

Approving isn't the same as judging.

SPLIT THE WORK

The work	AI	PEOPLE
Thinking	Suggests, challenges	Decide what matters
Doing	Drafts, sorts	Check, finish
Deciding	Never alone	Make the call

YOU'LL KNOW IT'S WORKING WHEN

- ✓ The calls that stay human are agreed in advance, each with a named owner.
- ✓ Training is on real tasks: when to use AI, how to check it, what to keep.
- ✓ Named champions have time set aside, share what works and report what breaks.

WATCH OUT FOR

- Signing off means approving whatever the tool suggested.

FIRST MOVE

- Before anything goes out, ask: could I defend this without saying "the AI suggested it"?

06 Prove it with a number

Agree the number before you build.

A demo shows AI at its best. A number shows an ordinary day.

GOOD FIRST NUMBERS

- 1 **Hours a week**
→ spent on the task
- 2 **Error rate**
→ in graded samples
- 3 **Turnaround**
→ from request to done

YOU'LL KNOW IT'S WORKING WHEN

- ✓ A baseline and a target are agreed before the build, with whoever judges success.
- ✓ A named owner tracks the number against the baseline.
- ✓ Every pilot has a decision date: scale, change or stop, on the number.

WATCH OUT FOR

- Success is reported in logins and licenses.

FIRST MOVE

- Measure one number before the next project starts.

Responsible AI: six rules

For every AI tool.

- 1 Write down what data may go into AI tools, and what never may.
- 2 Anything acting on customer data, or sending it outside, gets security review before it's built.
- 3 Anything that sends, pays, deletes or publishes gets a person in the loop or a hard limit.
- 4 Every AI tool has a named owner and a line on the list.
- 5 A named person makes every call that matters, and can defend it without the AI.
- 6 When AI decides about people, check its errors group by group: an average can hide who it fails.

The AI champion network

Champions are sensors, not cheerleaders.

- One per team, with a few hours a month for it.
- Every month they ask two questions: what's working that others should copy? What's breaking that someone should fix?

Where do you stand?

A free five-minute assessment scores all six plays and shows where to start. georgeandersen.net/tools/how-boring-is-your-ai

Who made this

George Andersen helps lead AI enablement for a large commercial organization. These plays come from that work.

linkedin.com/in/cgeorgeandersen

Built on five frameworks: [AI Should Be Boring](#), [Risk-Tiered AI Adoption](#), [Fix First, AI Last](#), [Calibrated Trust](#), [Keep It Human](#).

THE BORING AI PLAYBOOK · GEORGE ANDERSEN · [GEORGEANDERSEN.NET/PLAYBOOK](https://georgeandersen.net/playbook) · UPDATED OCT 2026 · FREE TO SHARE, WITH CREDIT.