


Responsible AI at CONA Services



Using AI well, using it safely, and proving it's worth it

Company-wide training · Built on Microsoft's Responsible AI Standard and the Cloud Adoption Framework for AI

TODAY

Five short modules, about 75 minutes

01	AI basics What it's good at, what it isn't, and how Copilot works	10 min
02	Microsoft's six principles The rules we use to judge any use of AI	15 min
03	Using AI safely every day Data, sharing, checking, and when to ask	15 min
04	How we govern AI projects Risk tiers, approvals and the paperwork	20 min
05	Proving the value How we measure what AI is worth	10 min

Plus a five-question knowledge check and your first 30 days.

Four things you'll be able to do



Use AI on real work

Brief Copilot like a colleague and get a useful first draft, not a generic one.



Apply the six principles

Ask the right questions before AI touches customers, people or sensitive data.



Know when it's a project

Spot when an idea needs a risk tier, a review and an approval.



Show what it's worth

Log wins and track benefits honestly, so the numbers stand up.

01

AI basics

What it's good at, what it isn't, and how Microsoft 365 Copilot works

A smart new hire who has never worked here

Well read, fast, eager to help. It doesn't know our customers, systems or last week's decision unless you tell it.



Good at

- Drafting emails, documents and summaries
- Rewriting: shorter, clearer, a different tone
- Pulling structure out of messy notes
- Explaining formulas, code and jargon
- Comparing two documents
- Brainstorming options and questions



Not good at

- Exact arithmetic: it can add up wrong and sound sure
- Facts it wasn't given: it may invent a policy or a statistic
- Knowing our context: customers, history, decisions
- Giving the same answer twice
- Judgement calls about people
- Knowing when it's wrong

Use the right tool for the job

A lot of work doesn't need AI. Knowing that is part of using it well.

The job	Best tool	Why
Summarise a 40-message email thread	Copilot in Outlook	Reading and condensing language is its strength
Total a column of 400 numbers	Excel	Exact every time. AI can explain the formula
Apply the same written rules to every request	Power Automate or a workflow	Rules run identically and leave an audit trail
Draft a first version of a proposal	Copilot in Word	A strong first draft saves hours; you edit
Decide who gets a pay rise	A person	Decisions about people stay with people
Report on something a system already tracks	The existing report	It's already built, supported and trusted

How Microsoft 365 Copilot works

It answers using the content you already have permission to see, inside our Microsoft 365 tenant.




You ask

A prompt in Word, Outlook, Teams, Excel or Copilot Chat



It looks

Searches the emails, chats, meetings and files you can already open



It drafts

A large language model writes the answer from what it found



You check

Answers show references. Open them before you rely on it

It respects permissions. Copilot never gives you access to anything you couldn't already open.

Your data isn't used to train the models. Microsoft states that prompts, responses and data accessed through Microsoft 365 aren't used to train its foundation models.

Brief it the way you'd brief a colleague

1

Who it's for

“For the VP of Sales, reading on their phone”

2

What to do

One clear verb: draft, shorten, compare, list

3

The material

Paste it or reference the file. Don't make it guess

4

What good looks like

Length, format, tone: “under 120 words, three bullets”

5

An example

“Write it like this one” beats any description

WEAK

Write an email about the export delay.

STRONG

Write an email to our main contact at [customer]. Their data export scheduled for Tuesday will now run Thursday because a migration failed overnight. We've run their three most-used reports manually so they aren't blocked. Firm and factual, one apology, under 120 words. End by offering a call tomorrow.

02

Microsoft's six principles

The questions we ask about any use of AI, from a quick prompt to a major project

Six principles, one question each

Microsoft's wording, and the question it turns into for us.



Fairness

AI systems should treat all people fairly.

Could this treat someone differently because of who they are?



Reliability and safety

AI systems should perform reliably and safely.

What happens when it's wrong, and would we catch it?



Privacy and security

AI systems should be secure and respect privacy.

Is this data allowed in this tool, and is it shared only as needed?



Inclusiveness

AI systems should empower everyone and engage all people.

Does it work for everyone who has to use it?



Transparency

AI systems should be understandable.

Do people know AI was involved, and what it can't do?



Accountability

People should be accountable for AI systems.

Who owns this, and who answers for it?



PRINCIPLE 1 OF 6

Fairness

"AI systems should treat all people fairly."

Microsoft Responsible AI Standard

The question: Could this treat someone differently because of who they are?

IN YOUR DAY-TO-DAY WORK

- Don't use AI to rank, screen or judge people.
- Watch for AI copying biased language from old material.
- If an output treats groups differently, flag it.

IN AI PROJECTS

- Is it a sensitive use: jobs, pay, credit, access to services?
- Test outputs across different customer and employee groups.
- A person makes every decision about people.

EXAMPLE

Using AI to shortlist job applicants is a sensitive use and needs a Tier 3 review. Using it to tidy your own interview notes is fine.



PRINCIPLE 2 OF 6

Reliability and safety

“AI systems should perform reliably and safely.”

Microsoft Responsible AI Standard

The question: What happens when it's wrong, and would we catch it?

IN YOUR DAY-TO-DAY WORK

- Check numbers, names and dates against the source.
- Don't rely on AI for anything you can't check.
- Use Excel or a report for exact figures.

IN AI PROJECTS

- Test on real past examples before launch.
- Set the pass mark before the pilot starts.
- Monitor it once it's live, and have a way to switch it off.

EXAMPLE

A support drafting pilot was tested on 50 past tickets, with a pass mark written down before it started, and a feature flag to turn it off in minutes.



PRINCIPLE 3 OF 6

Privacy and security

“AI systems should be secure and respect privacy.”

Microsoft Responsible AI Standard

The question: Is this data allowed in this tool, and is it shared only as needed?

IN YOUR DAY-TO-DAY WORK

- Use approved tools only, like Copilot in our tenant.
- Don't paste personal or confidential data you don't need.
- Keep sensitivity labels on, and share on purpose.

IN AI PROJECTS

- Privacy review for personal data.
- Security review for any new vendor or model.
- Only the data fields that are actually needed.

EXAMPLE

Replacing names and account numbers with placeholders like [Customer A] lets AI help without the personal data ever leaving your hands.



PRINCIPLE 4 OF 6

Inclusiveness

“AI systems should empower everyone and engage all people.”

Microsoft Responsible AI Standard

The question: Does it work for everyone who has to use it?

IN YOUR DAY-TO-DAY WORK

- Write prompts and outputs in plain language.
- Consider people using screen readers or a second language.
- Don't leave out colleagues who work differently.

IN AI PROJECTS

- Who uses it, and who is affected by it?
- Test accessibility and language support.
- Keep a non-AI route for people who need it.

EXAMPLE

A drafting tool rolled out in English and Spanish, tested with a screen reader, and kept the old workflow for agents who opted out.



PRINCIPLE 5 OF 6

Transparency

“AI systems should be understandable.”

Microsoft Responsible AI Standard

The question: Do people know AI was involved, and what it can't do?

IN YOUR DAY-TO-DAY WORK

- Don't present AI output as a checked fact.
- Say when AI drafted something, where our policy asks you to.
- Open the references Copilot shows.

IN AI PROJECTS

- Tell users when they're dealing with AI output.
- Explain its limits in plain words.
- Customers get replies from a named person.

EXAMPLE

Agents see an “AI draft” label and a one-page guide to its limits. Customers still get replies written and sent by a named agent.



PRINCIPLE 6 OF 6

Accountability

"People should be accountable for AI systems."

Microsoft Responsible AI Standard

The question: Who owns this, and who answers for it?

IN YOUR DAY-TO-DAY WORK

- You own everything you send, whoever drafted it.
- "The AI did it" is never the answer.
- Report problems to the tool owner.

IN AI PROJECTS

- A named owner for every AI project.
- A named reviewer for outputs, and a route for complaints.
- Approvals by people who aren't involved in the project.

EXAMPLE

Owner: Head of Support Operations. Agents own every reply they send. Issues reach the support ops lead within one business day.

Three sensitive uses need the highest level of review

From Microsoft's Responsible AI Standard. Any “yes” makes a project Tier 3 in our process.



Life opportunities

Could it affect someone's legal position or life opportunities, such as employment, pay, credit, housing, insurance, education or access to services?



Serious harm

Could its use or misuse cause someone significant physical or psychological injury?



Human rights

Could it restrict or undermine someone's human rights, such as privacy, freedom of expression or equal treatment?

Not sure? Treat it as a yes and ask the AI review board. A sensitive use isn't a no. It's a yes with the right people in the room.

QUICK CHECK

Which principle is most at stake?

Discuss in pairs for two minutes. Answers are in the speaker notes.

A

A chatbot answers customers without anyone checking, and nobody told them it's AI.

B

Copilot surfaces a salary spreadsheet shared with the whole company by mistake.

C

An AI tool ranks employees for the annual performance review.

D

A report generator gets figures wrong and nobody notices for a month.

03

Using AI safely every day

Approved tools, careful sharing, and checking before anything leaves your hands

What's safe to use with AI

Our AI policy is the final word. These are the common-sense basics.



Go ahead

- General business information in approved tools
- Your own notes, drafts and documents
- Public information



Take care

- Personal data about customers or employees
- Confidential or labelled documents
- Only in approved tools, and only the fields you need



Don't

- Company data in personal or unapproved AI tools
- Passwords, keys or payment card data
- Regulated or NDA data without written approval

EVERY DAY

Copilot sees what you can see, so share on purpose

It respects your permissions. Copilot only uses content you can already open.

But it finds things fast. A folder shared with everyone by mistake can show up in anyone's answer.

Sensitivity labels matter. Label confidential files. Copilot respects labels, and content it creates from a labelled file picks up the label.

Check the references. If Copilot cites something you didn't expect to see, tell the owner.

RISKY

Salary review spreadsheet shared with “People in your organization” so two managers could open it.

Now it can appear in anyone's Copilot answer about pay.

SAFE

Shared with the two managers by name, with the Confidential sensitivity label applied.

They can open it. Nobody else's Copilot will surface it.

Check it before it leaves your hands

AI writes fluently even when it's wrong. Your name goes on what you send.



Facts

Numbers, names, dates and quotes match the source. These are the most common errors.



Promises

Did it commit us to a date, a refund or a feature nobody agreed to?



Tone

Cut the over-apologising and padding. Would you say it out loud?



Gaps

Ask it: "What did you leave out?" Summaries quietly drop the awkward bit.

Be most careful where you know least. If you couldn't spot a wrong answer, ask someone who could.

Personal use, or a project?

Most AI use is personal productivity and needs no approval. Some needs a project.

Personal use: just do it well

- You're drafting, summarising or rewriting your own work
- A person reads it before it goes anywhere
- It uses data the tool is approved for
- Nobody else depends on it running

Start a project if any is true

- Output reaches customers or the public
- It affects decisions about people
- It uses personal, confidential or regulated data at scale
- It acts in other systems automatically
- It needs a new vendor, model or tool
- A team or process will depend on it

04

How we govern AI projects

Following the Govern AI steps in Microsoft's Cloud Adoption Framework

Microsoft's four steps for governing AI

From the Govern AI guidance in Microsoft's Cloud Adoption Framework, and how we apply each one.




1. Assess

Understand each AI use and its risks, using the six principles.

HOW WE DO IT

Intake, risk screen with sensitive uses, Responsible AI assessment




2. Document

Write down the policies that address those risks.

HOW WE DO IT

Risk tiers, required paperwork, approval authority, training




3. Enforce

Apply the policies, automated where possible.

HOW WE DO IT

Stage gates, board-only independent approvals, Purview labels



4. Monitor


Keep checking risk and results once it's running.

HOW WE DO IT

Quarterly or yearly risk reviews, benefits tracking, post-launch review

Every project follows the same seven stages

A project can't move on until the paperwork for its stage is done.



Gates, not suggestions. The portal won't move a project to the next stage until the required paperwork is complete. Only the AI review board can record an approval, and never on a project they started.

Three risk tiers decide how much review a project needs

Tier 1

Low risk

Internal use, general data, and a person checks the output.

APPROVED BY

Team lead and AI governance lead

RISK REVIEW ONCE LIVE

Every year

Tier 2

Medium risk

Personal or confidential data, automated actions, finance processes, or a new vendor.

APPROVED BY

AI review board

RISK REVIEW ONCE LIVE

Every year

Tier 3

High risk

A Microsoft sensitive use, or output reaching customers without a person checking it.

APPROVED BY

AI review board with Legal and Security

RISK REVIEW ONCE LIVE

Every quarter

The highest “yes” in the risk screen sets the tier.

The paperwork each tier needs

All of it lives in the project's page in the AI portal, and exports as one governance pack PDF.

Paperwork	What it answers	T1	T2	T3
Intake and idea check	Is this a job for AI at all?	●	●	●
Risk tier	How risky is it, using the sensitive uses and screening?	●	●	●
Exec brief	What's the ask, the problem, the plan and the numbers?	●	●	●
Approval record	Who decided, when, and on what conditions?	●	●	●
Pilot plan and result	What pass mark did we set, and did we hit it?		●	●
Responsible AI assessment	How does it meet each of the six principles?		●	●
Privacy and legal sign-off	Have Privacy and Legal reviewed it before launch?			●
Benefits tracking	What did it actually deliver, quarter by quarter?	●	●	●
Post-launch review	What worked, what didn't, keep or stop?	●	●	●

The Responsible AI assessment

One answer per principle. Questions adapted from Microsoft's Cloud Adoption Framework. Required for Tier 2 and 3.

Principle	Microsoft's question	What to write
Fairness	How could it lead to unequal treatment or bias?	Who could be treated differently, and how you'll check
Reliability and safety	Where could it fail or give unreliable results?	Testing, monitoring and how to switch it off
Privacy and security	How might it expose sensitive data or create a security gap?	What data it touches, where it goes, who approved the tool
Inclusiveness	Could some groups be left out or disadvantaged?	Who uses it or is affected, and how it works for them
Transparency	What would be hard for people to understand?	Who knows AI is involved and how its limits are explained
Accountability	Where could responsibility be unclear?	The named owner, reviewer and complaints route

Who does what



Everyone

Uses approved tools, checks outputs, logs wins, and flags anything that feels wrong.



Executive sponsor

Backs the project, removes blockers, and is accountable for its business case.



Legal, Privacy, Security

Sign off Tier 3 projects, review new vendors, and advise on sensitive data.



Project owner

Runs the project, completes the paperwork, tracks benefits, and owns the outcome.



AI review board

Records approval decisions. Never approves a project its members started.



AI governance lead

Keeps the process running, approves Tier 1 with team leads, and reports to leadership.

Going live isn't the end: monitor it

Microsoft's guidance: keep evaluating risk, measure consistently, and report what you find.



Scheduled risk reviews

Every quarter for Tier 3, every year otherwise. Overdue reviews show in red on the portal.



Benefits tracking

Planned against actual, every quarter, with the evidence for each number.



Monitoring

The measures named in the Responsible AI assessment, reviewed by the named person.



Switch-off plan

Who can turn it off, and how fast, if something goes wrong.



Post-launch review

About three months in: what worked, what didn't, and whether to scale, change or stop.

05

Proving the value

Honest numbers that stand up in front of leadership

Hours saved are not money saved

Time freed up only becomes value when it goes somewhere.

Counts

- A hire we no longer need
- Overtime or contractors that stop
- Named work people move to

Doesn't count

- Time that just gets absorbed
- “We'll find something for them to do”
- Savings nobody can point to

Confidence haircut

100%	High: measured it
75%	Medium: trial or good estimate
50%	Low: educated guess

A benefit you're 50/50 on is worth about half of itself today. Costs never get a haircut.

Three numbers leadership will see every quarter

1

Adoption

People trained and active, by team

2

Time back


Wins people log, counted at 50% because self-reports run high

3

Delivered value

Project benefits against plan, with evidence

Planned vs delivered, \$k to date (illustrative)



Five questions. Answers in the speaker notes.

- 1 You need the exact total of a 400-row column. What do you use?
- 2 Copilot shows you a document you weren't meant to see. What most likely happened, and what do you do?
- 3 Name Microsoft's three sensitive uses.
- 4 A pilot has finished. When should its pass mark have been set?
- 5 A project saves 2,000 hours a year, but nobody has said where the time goes. How much of that counts?

WHAT NEXT

Your first 30 days

Everything below is in the CONA AI Portal.

WEEK 1



Finish the eight lessons

About 40 minutes. Takes you to Practitioner.

WEEK 2



Pick one weekly task

Use a prompt from the library on it every week.


WEEK 3



Log your wins

Twenty seconds each. It's how we prove this works.

WEEK 4



Share what works

Post a prompt, or run the 30-minute team session.

Levels: Starter → Practitioner → Regular → Champion

Where to go for help

CONA AI Portal

Lessons, prompts, projects and the value dashboard.

[\[Add the portal link\]](#)

AI review board

Questions about projects, tiers and approvals.

[\[Add a name or Teams channel\]](#)

IT and Security

Approved tools, sharing settings and sensitivity labels.

[\[Add the service desk link\]](#)

Microsoft's Responsible AI principles

The six principles and Microsoft's approach.

microsoft.com/ai/principles-and-approach

Cloud Adoption Framework: Govern AI

Microsoft's four steps for governing AI.

learn.microsoft.com/azure/cloud-adoption-framework/ai/govern

Microsoft Responsible AI Standard v2

Where the sensitive uses and impact assessment come from.

Search: Microsoft Responsible AI Standard General Requirements

**Use it well.
Check it.
Own it.**

Questions?

